

Adversarial search

EMI | Semestre 1 | Pr Mohamed RHAZZAF

A brief history

- **Checkers:**
 - 1950: First computer player
 - 1994: First computer world champion: Chinook defeats Tinsley
- **Chess:**
 - 1945-1960: Zuse, Wiener, Shannon, Turing, Newell & Simon, McCarthy
 - 1997: Deep Blue defeats human champion Gary Kasparov
- **Go:**
 - 1968: Zobrist's program plays legal Go, barely ($b>300!$)
 - 1968-2005: various ad hoc approaches tried, novice level
 - 2005-2014: Monte Carlo tree search -> strong amateur
 - 2016-2017: AlphaGo defeats human world champions

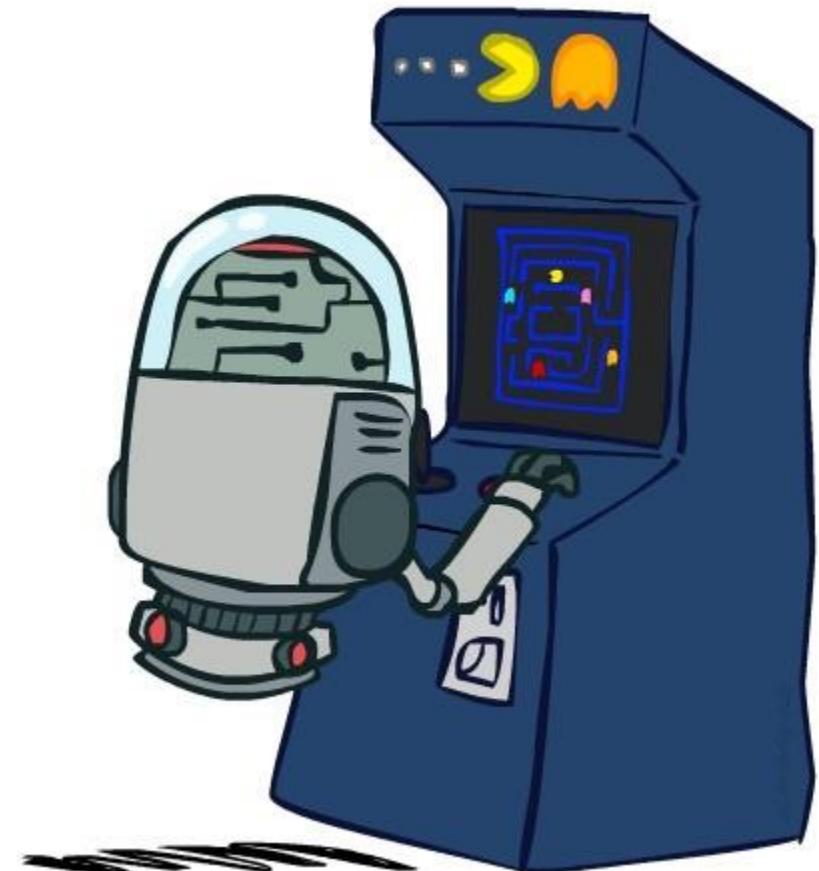
Types of Games

- Game = task environment with > 1 agent
- Axes:
 - Deterministic or stochastic?
 - Perfect information (fully observable)?
 - One, two, or more players?
 - Turn-taking or simultaneous?
 - Zero sum?
- Want algorithms for calculating a contingent plan (a.k.a. **strategy or policy**) which recommends a move for every possible eventuality



“Standard” Games

- Standard games are deterministic, observable, two-player, turn-taking, zero-sum
- Game formulation:
 - Initial state: s_0
 - Players: Player(s) indicates whose move it is
 - Actions: Actions(s) for player on move
 - Transition model: Result(s,a)
 - Terminal test: Terminal-Test(s)
 - Terminal values: Utility(s,p) for player p
 - Or just Utility(s) for player making the decision at root



Zero-Sum Games



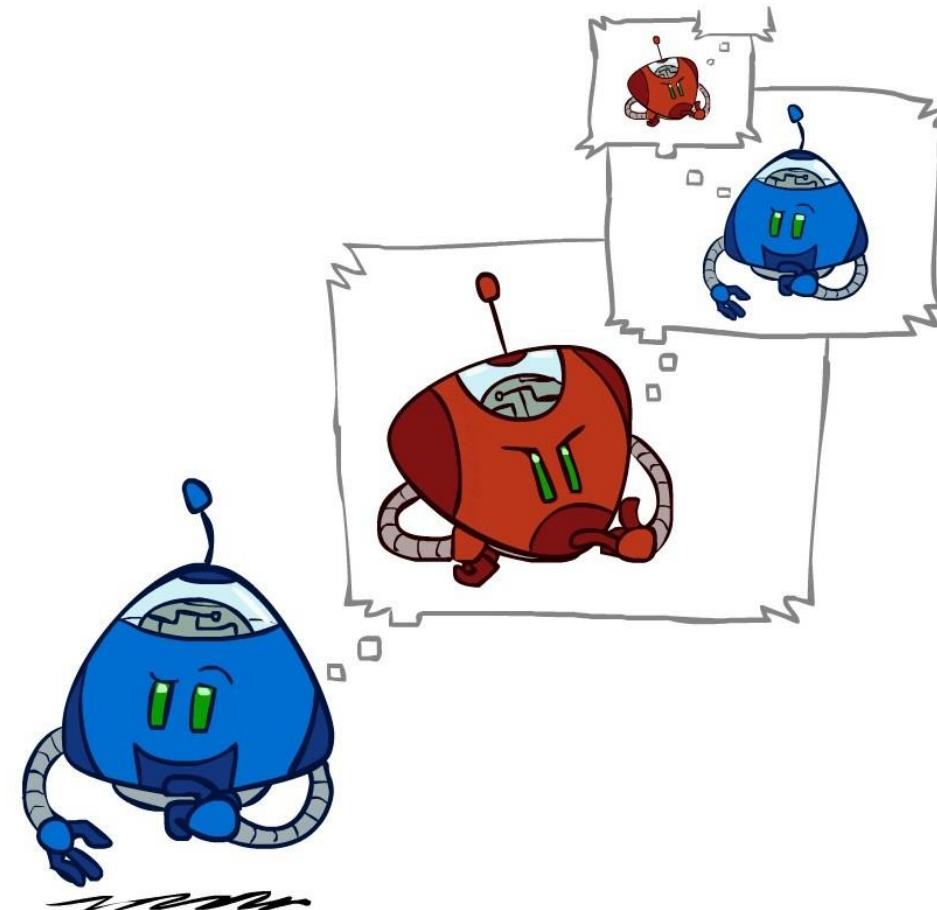
■ Zero-Sum Games

- Agents have **opposite utilities**
- Pure competition:
 - One **maximizes**, the other **minimizes**

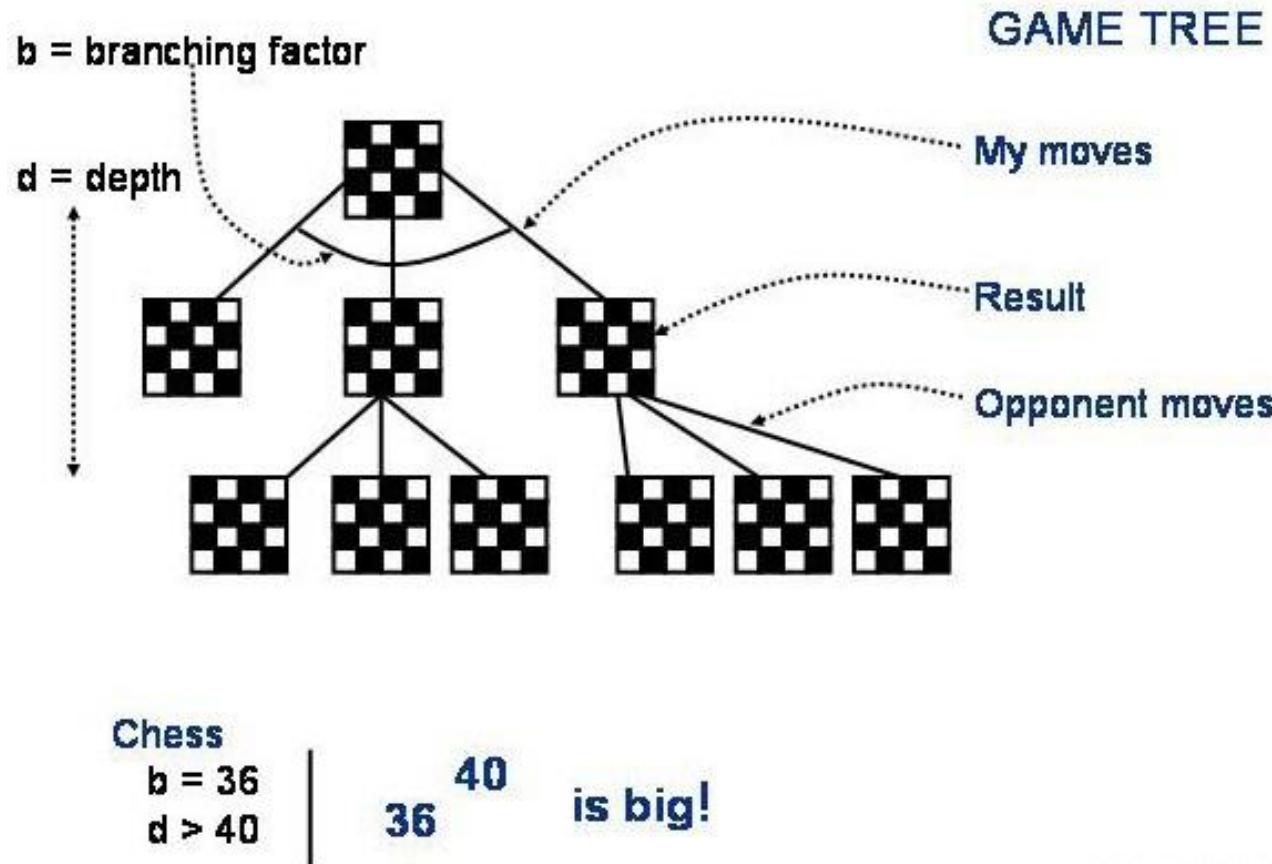
■ General Games

- Agents have independent utilities
- Cooperation, indifference, competition, shifting alliances, and more are all possible

Adversarial Search

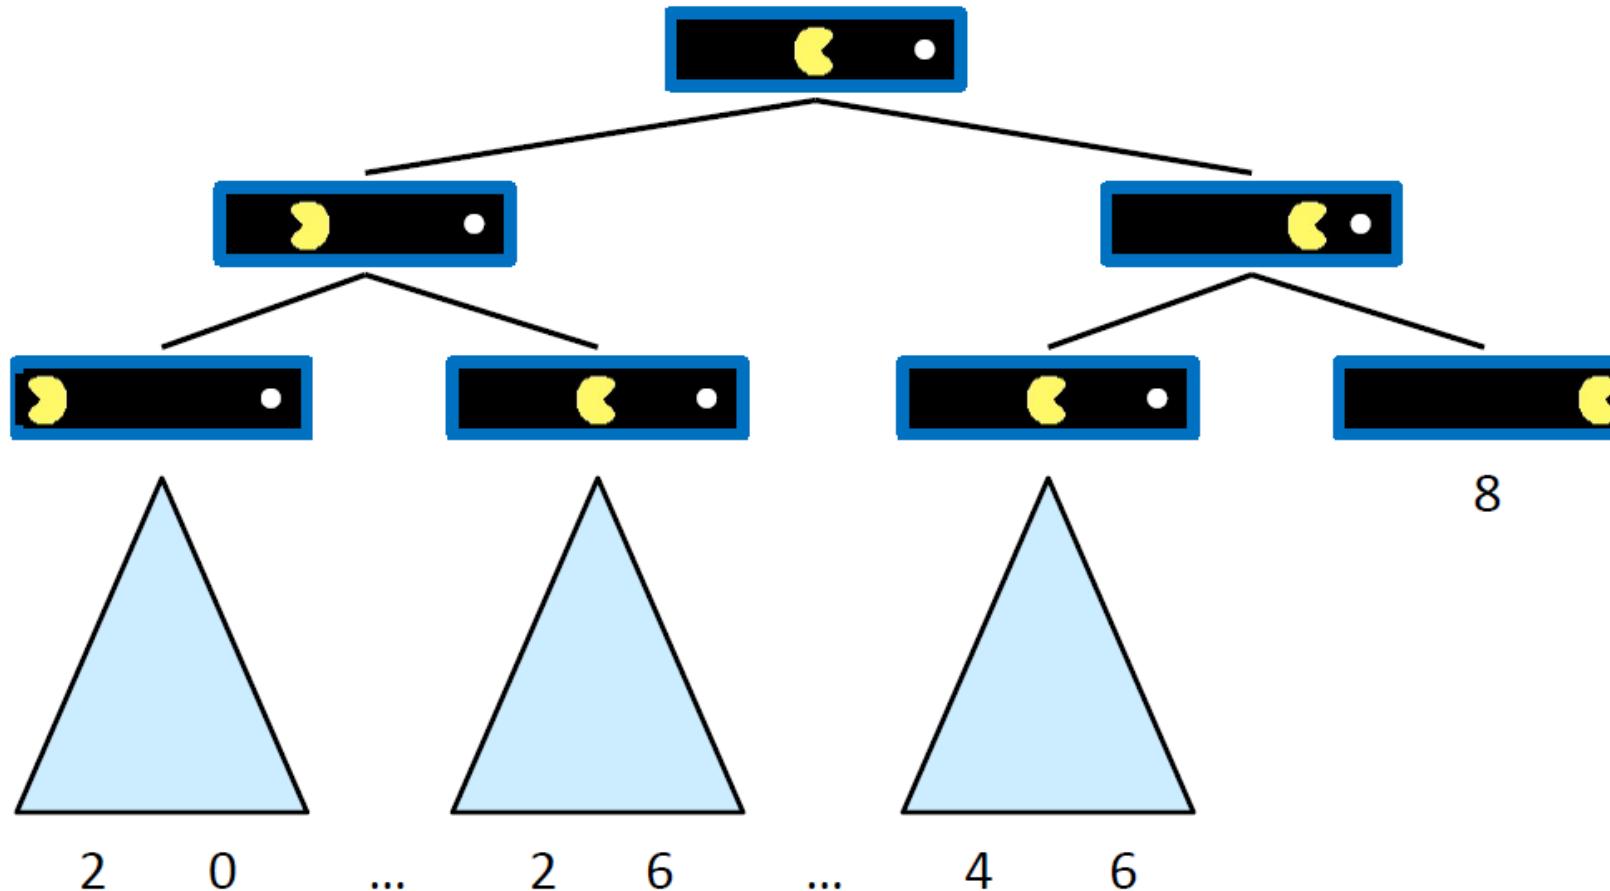


Two players games



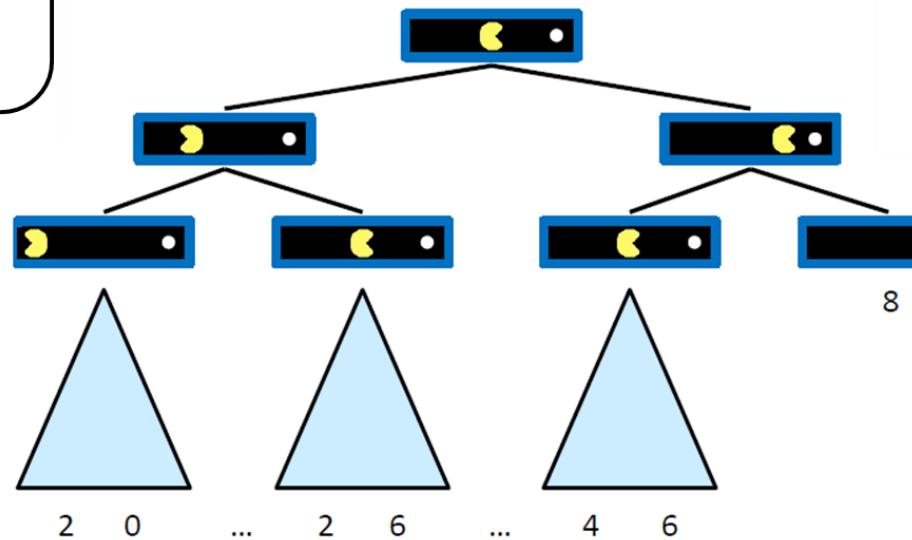
- Some board position represents the initial state
- We generate the children of this position by making all of the legal moves available to us
- Then, we consider the moves that our opponent can make to generate the descendants of each of these positions, etc.
- Note that these trees are enormous and cannot be explicitly represented in their entirety for any complex game.

Single-Agent Trees



Value of a State

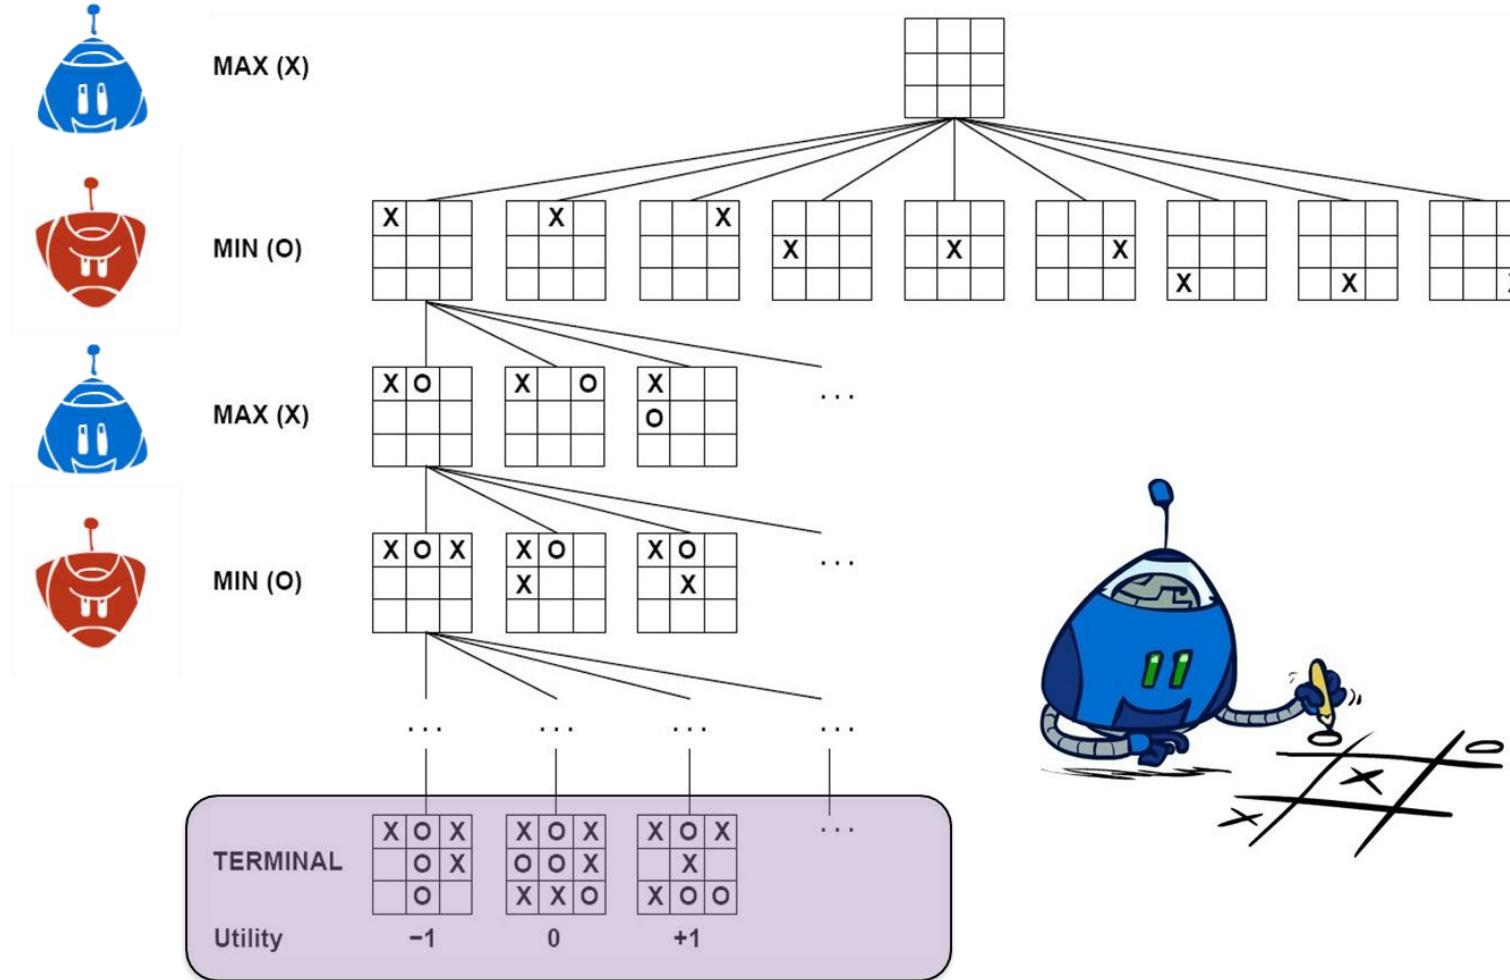
Value of a state:
The best achievable
outcome (utility)
from that state



Non-Terminal States:
$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

Terminal States:
$$V(s) = \text{known}$$

Tic-Tac-Toe Game Tree



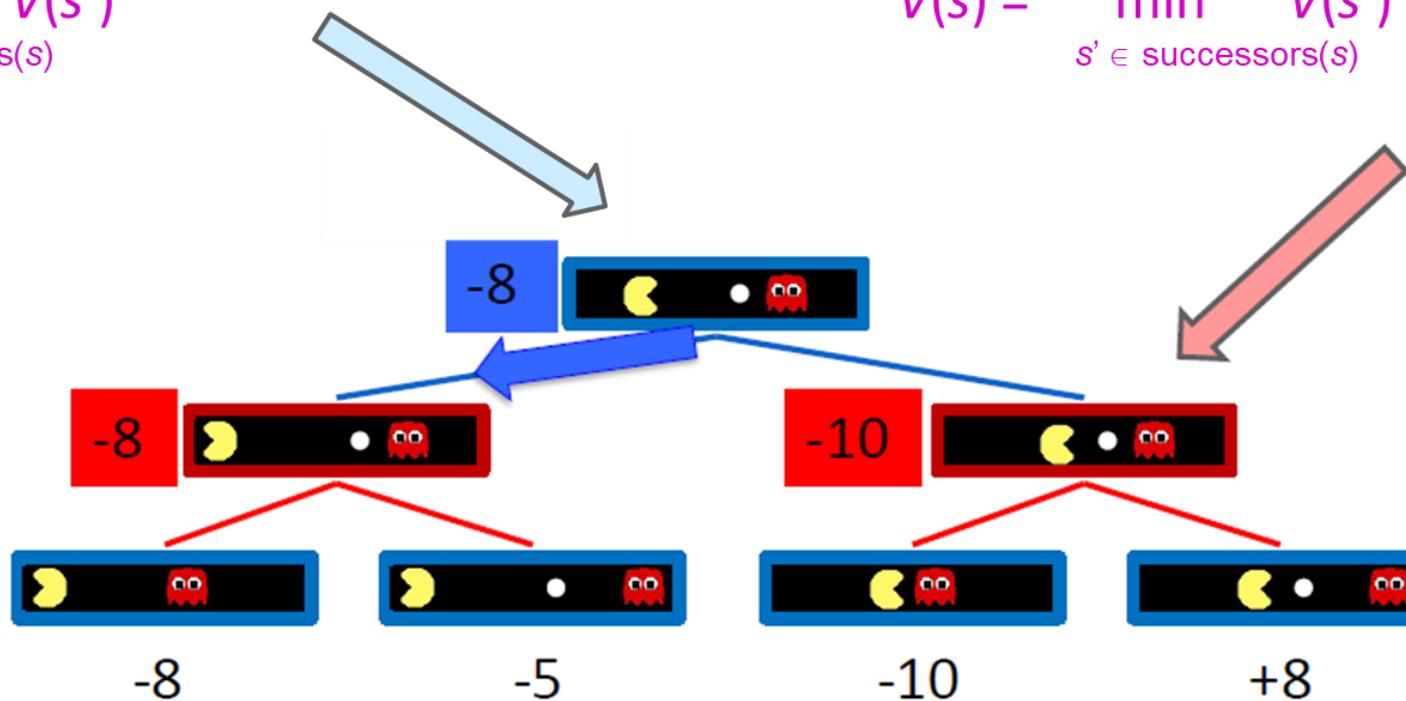
Minimax Values

MAX nodes: under Agent's control

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

MIN nodes: under Opponent's control

$$V(s) = \min_{s' \in \text{successors}(s)} V(s')$$



Terminal States:

$$V(s) = \text{known}$$

Minimax Algorithm

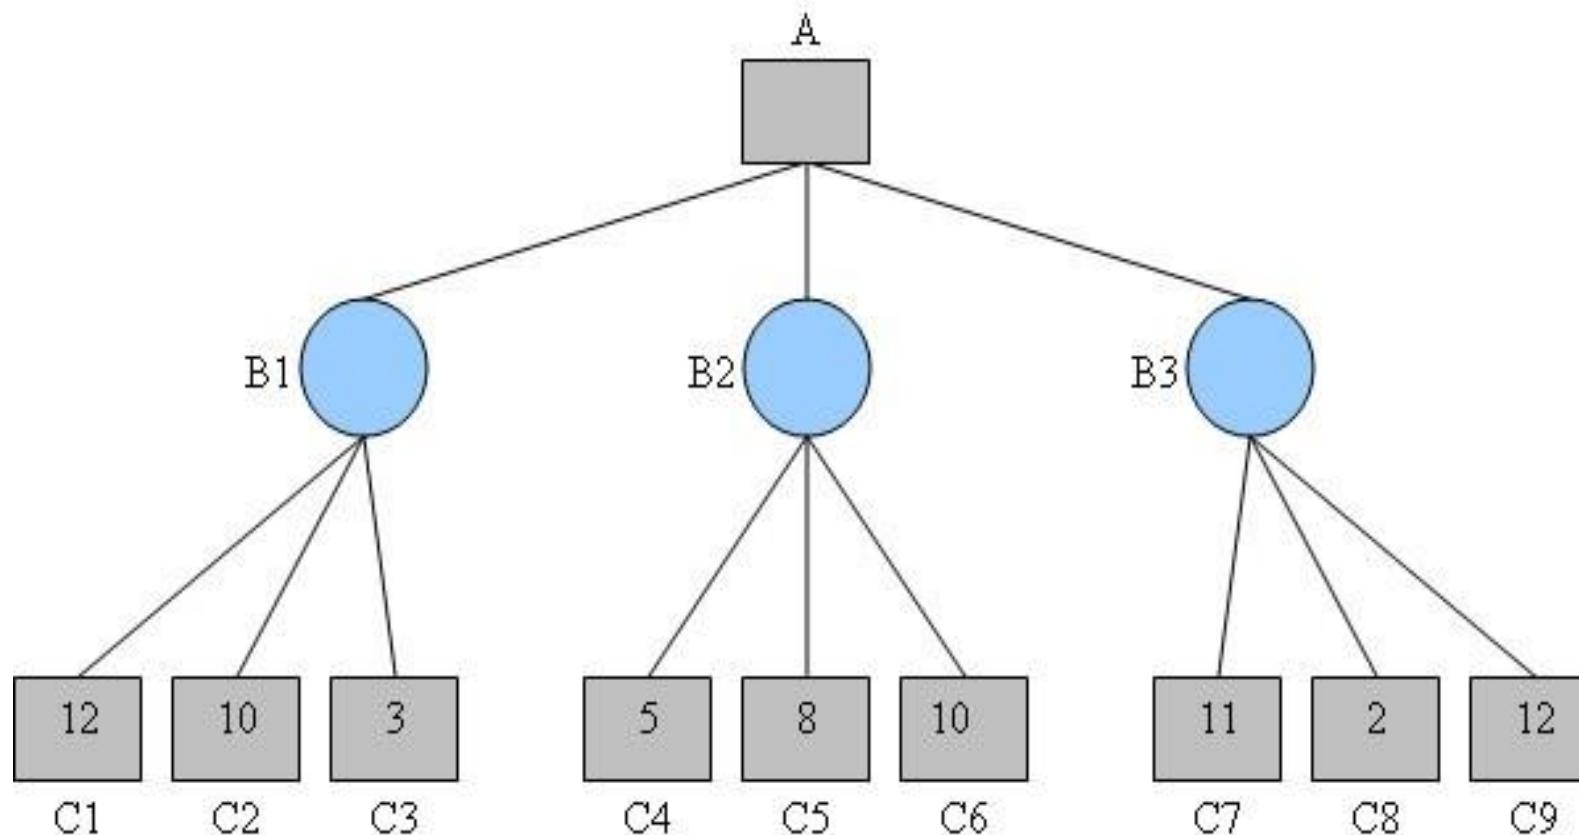
- **Problem setting**

- Two players:
 - **MAX**: tries to maximize the utility
 - **MIN**: tries to minimize the utility (because in zero-sum, MIN's gain = MAX's loss)
- Both players are assumed to play **optimally**.
- A game is represented by a **game tree**:
 - Nodes = states
 - Edges = actions
 - Leaves = terminal states with utilities for MAX

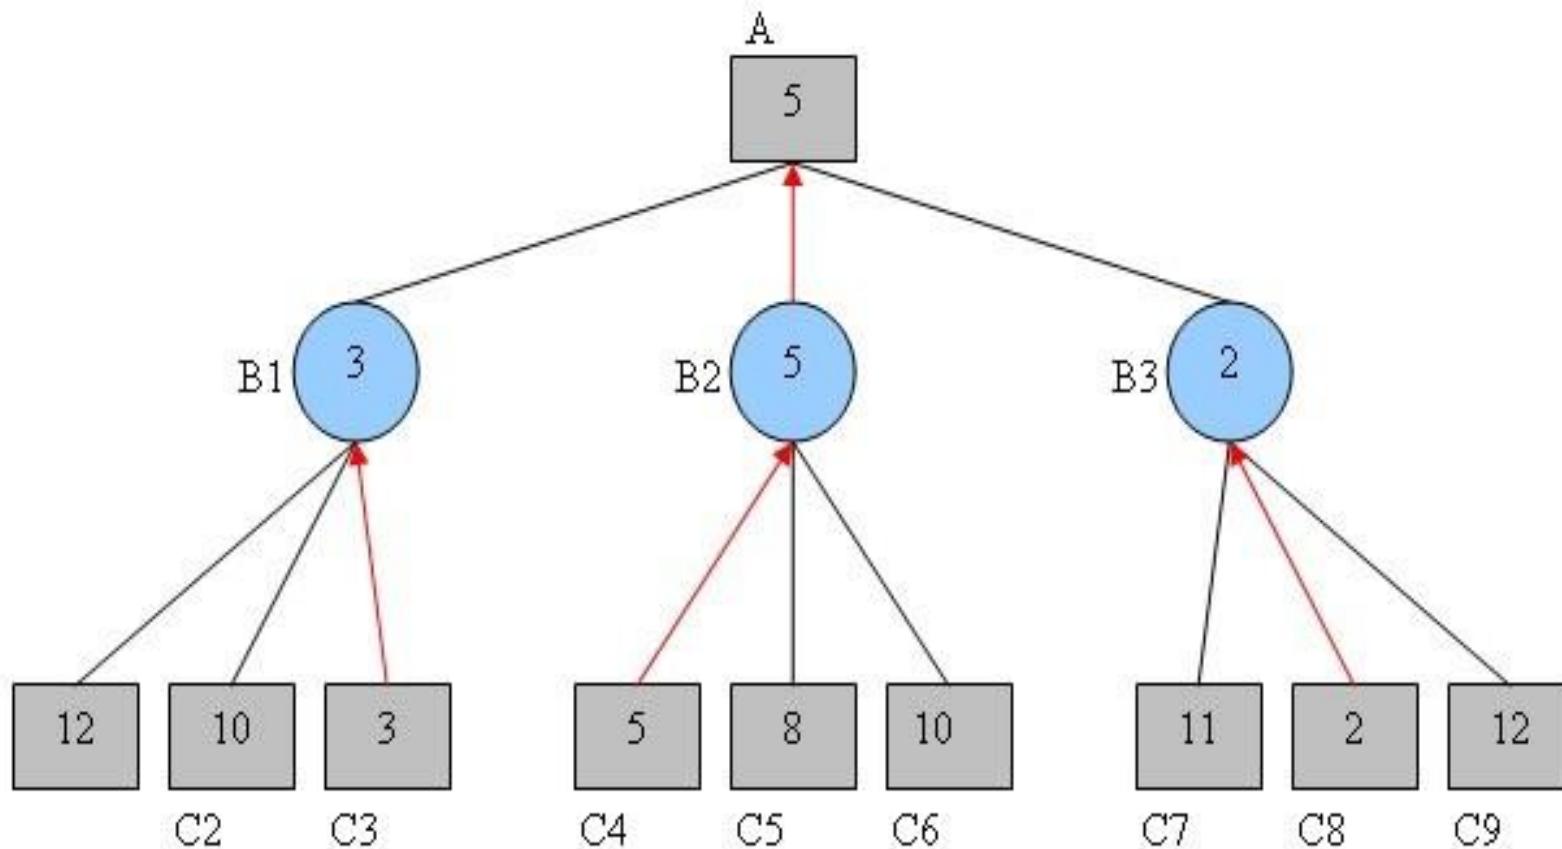
Minimax Algorithm

- Key Idea of Minimax
 - At each state:
 - MAX node chooses the action with the **maximum** value
 - MIN node chooses the action with the **minimum** value
 - Terminal nodes have fixed utility values (e.g., win = +1, loss = -1, draw = 0).
 - Values propagate **up the tree**.

Minimax Algorithm



Minimax Algorithm



Minimax Algorithm

```
def minimax_decision(state):
    """
    Returns the optimal action for MAX in the given state.
    """
    best_value = float("-inf")
    best_action = None

    for action in actions(state):
        value = minimax_value(result(state, action))
        if value > best_value:
            best_value = value
            best_action = action

    return best_action
```

```
def minimax_value(state):
    """
    Returns the minimax value of a given state.
    """

    # If the game is over, return its utility
    if terminal_test(state):
        return utility(state)

    # If it's MAX's turn
    if player(state) == "MAX":
        value = float("-inf")
        for action in actions(state):
            value = max(value, minimax_value(result(state,
action)))
        return value

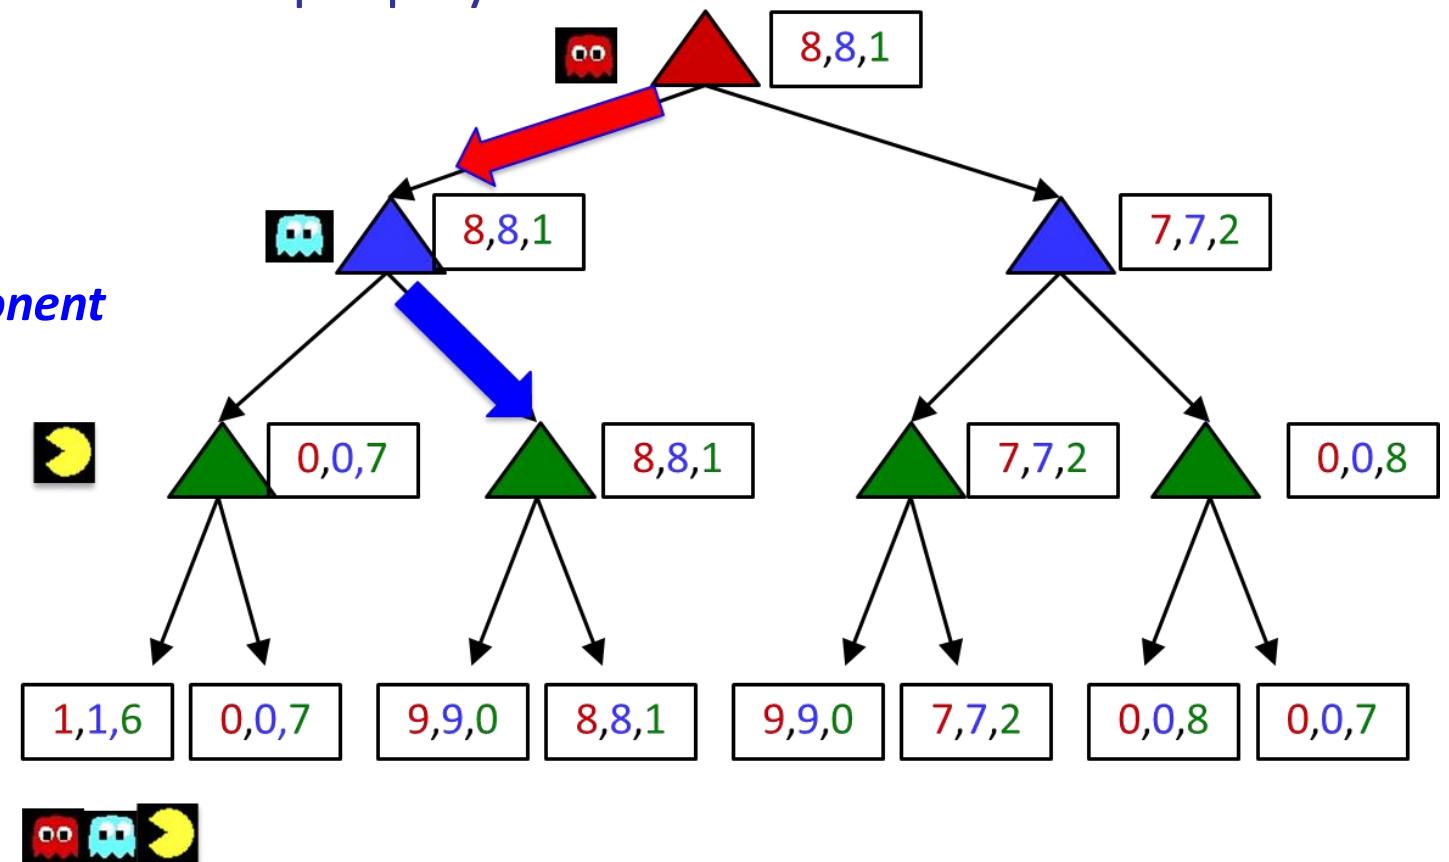
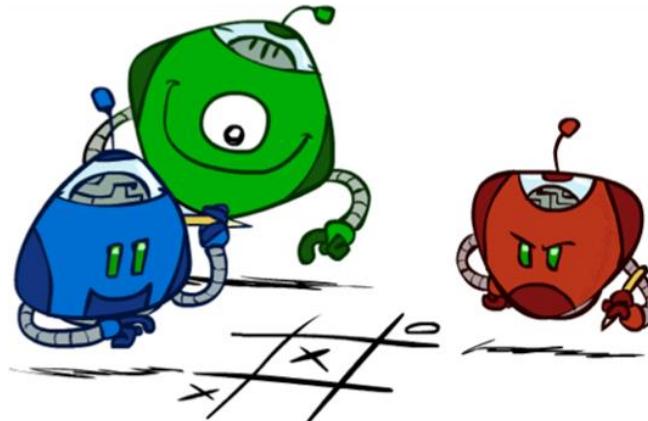
    # If it's MIN's turn
    else: # player(state) == "MIN"
        value = float("inf")
        for action in actions(state):
            value = min(value, minimax_value(result(state,
action)))
        return value
```

Generalized Minimax

- What if the game is not zero-sum, or has multiple players?

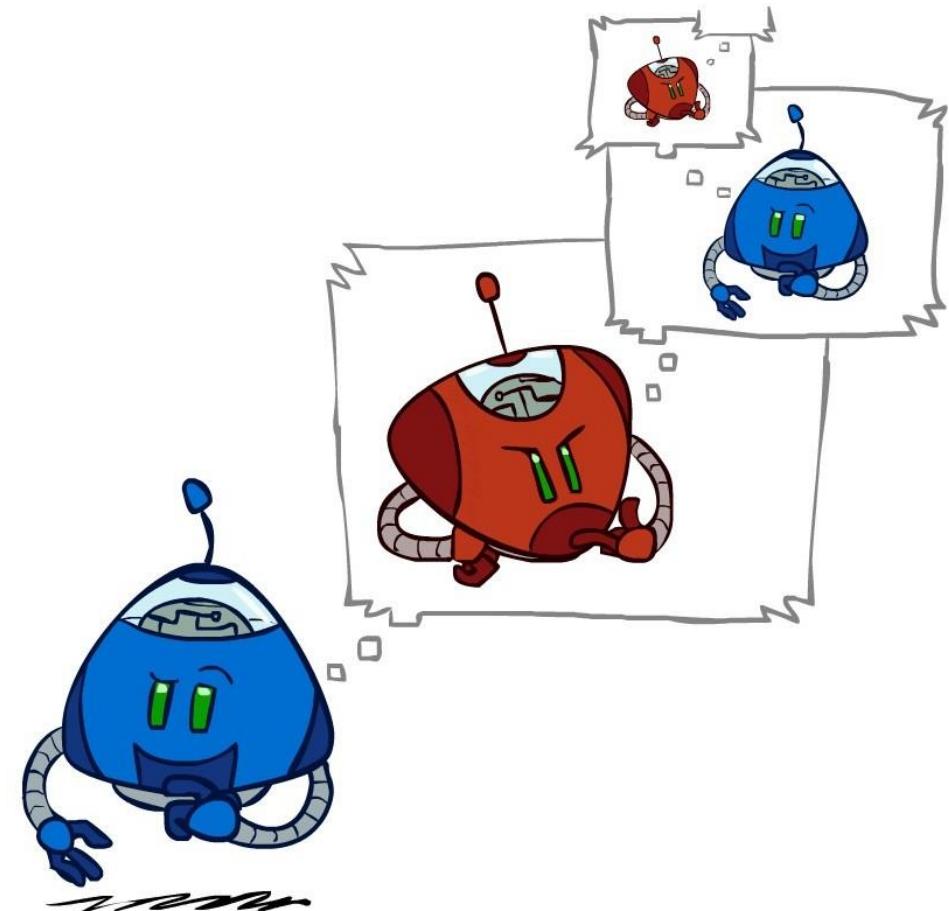
- Generalization of minimax:

- Terminals have **utility tuples**
- Node values are also utility tuples
- **Each player maximizes its own component**
- Can give rise to cooperation and competition dynamically...



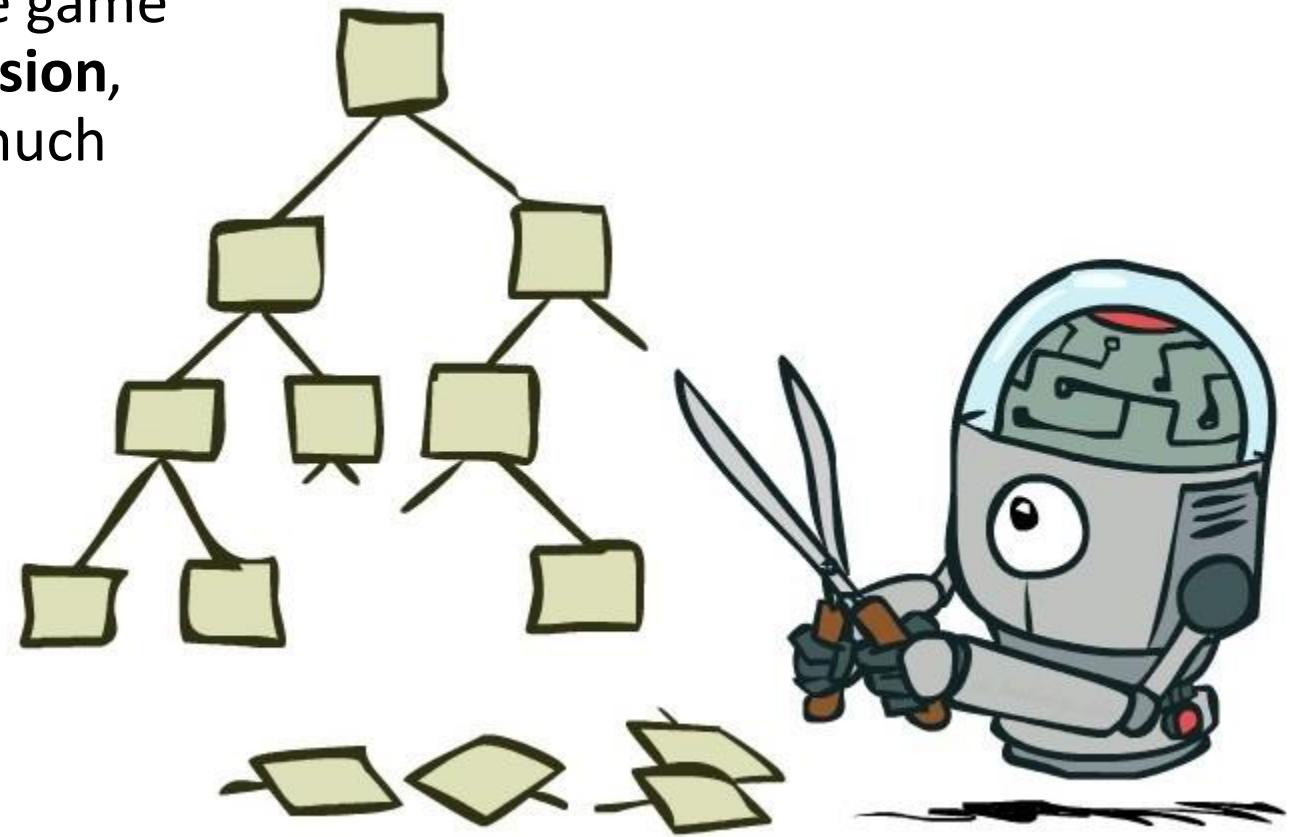
Minimax Efficiency

- **How efficient is minimax?**
 - Just like (exhaustive) DFS
 - Time: $O(b^m)$
 - Space: $O(bm)$
- **Example: For chess, $b = 35$, $m = 100$**
 - Exact solution is completely infeasible
 - Humans can't do this either, so how do we play chess?

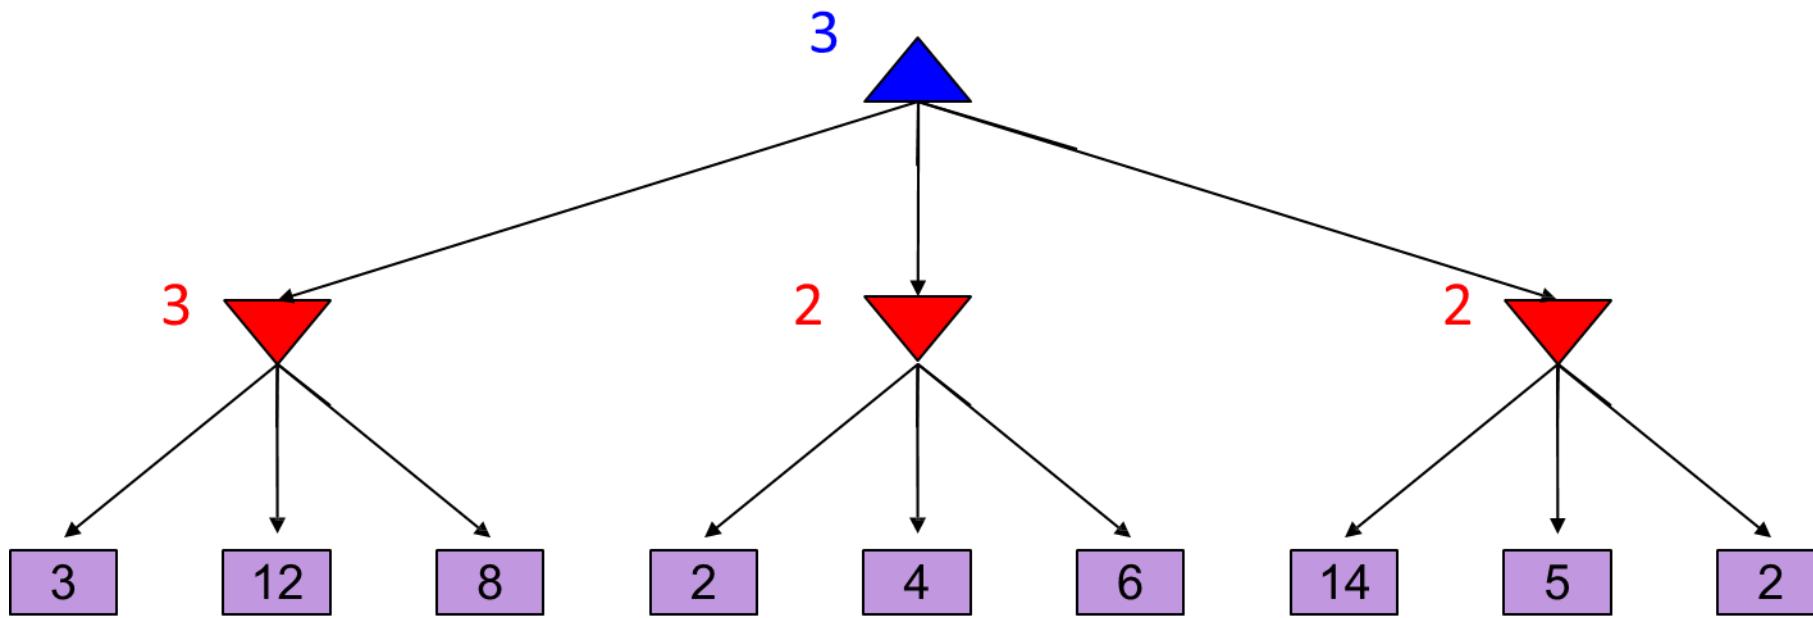


Game Tree Pruning

- **Pruning** refers to eliminating parts of the game tree that **cannot influence the final decision**, allowing the minimax algorithm to run much faster without changing the result

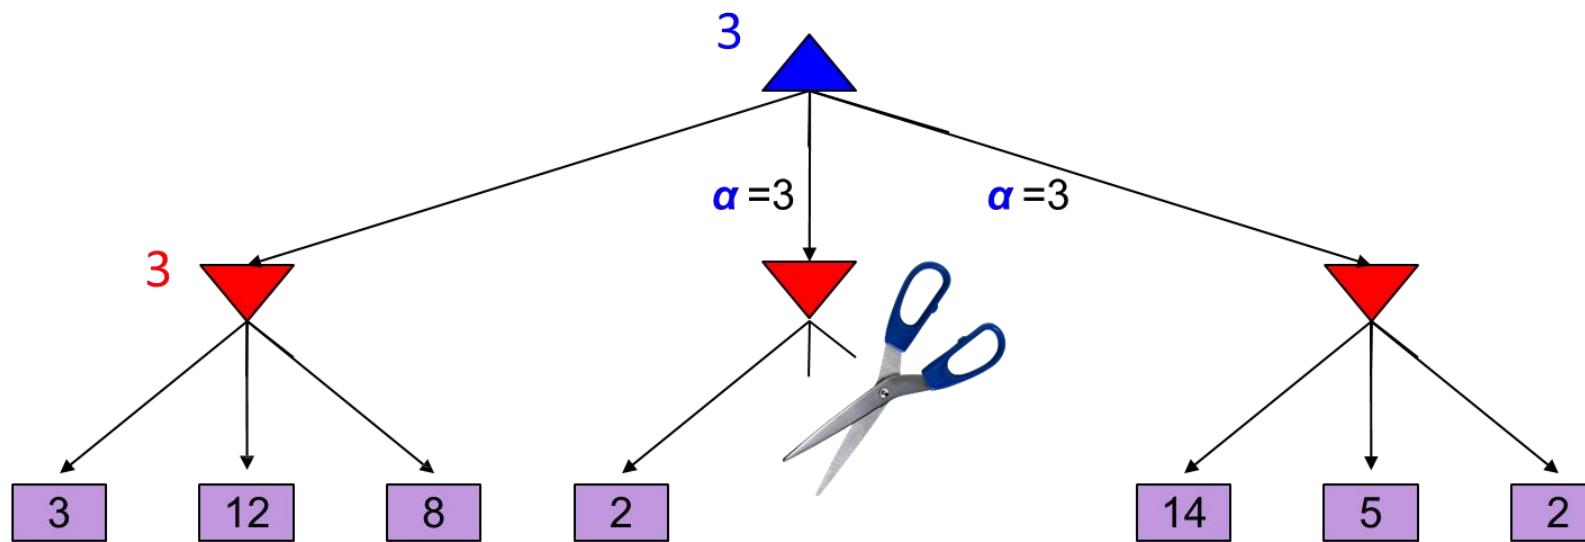


Minimax Example



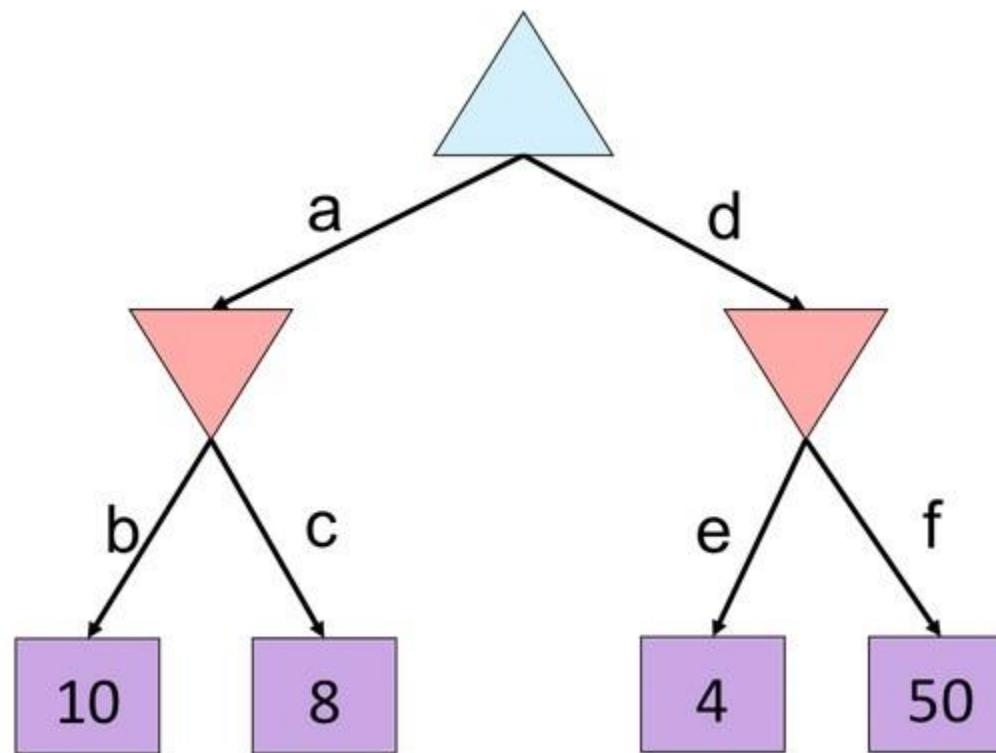
Alpha-Beta Example

α = best option so far from any MAX node on this path

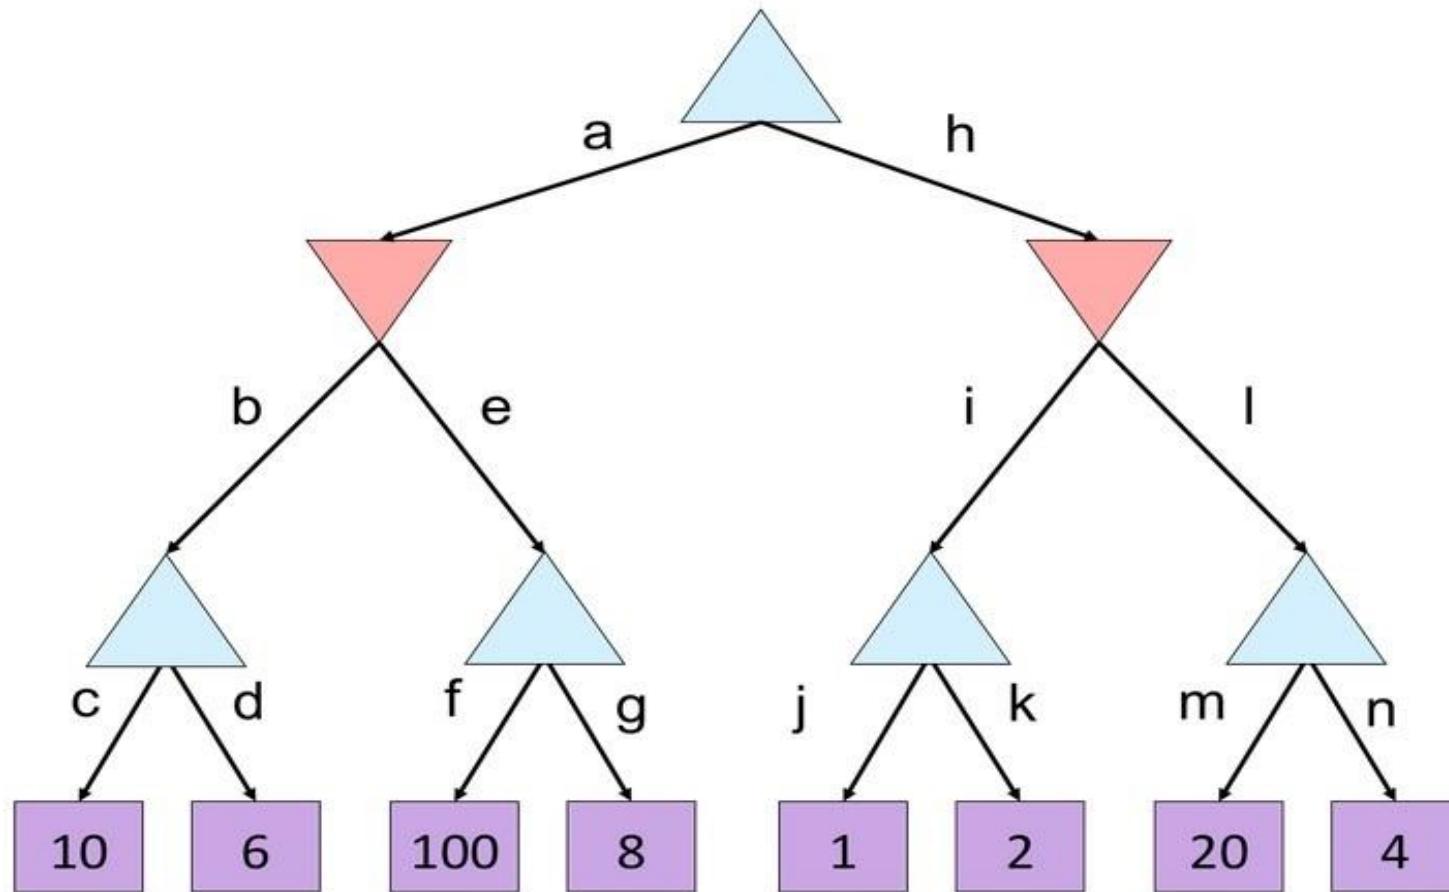


The order of generation matters: more pruning is possible if good moves come first

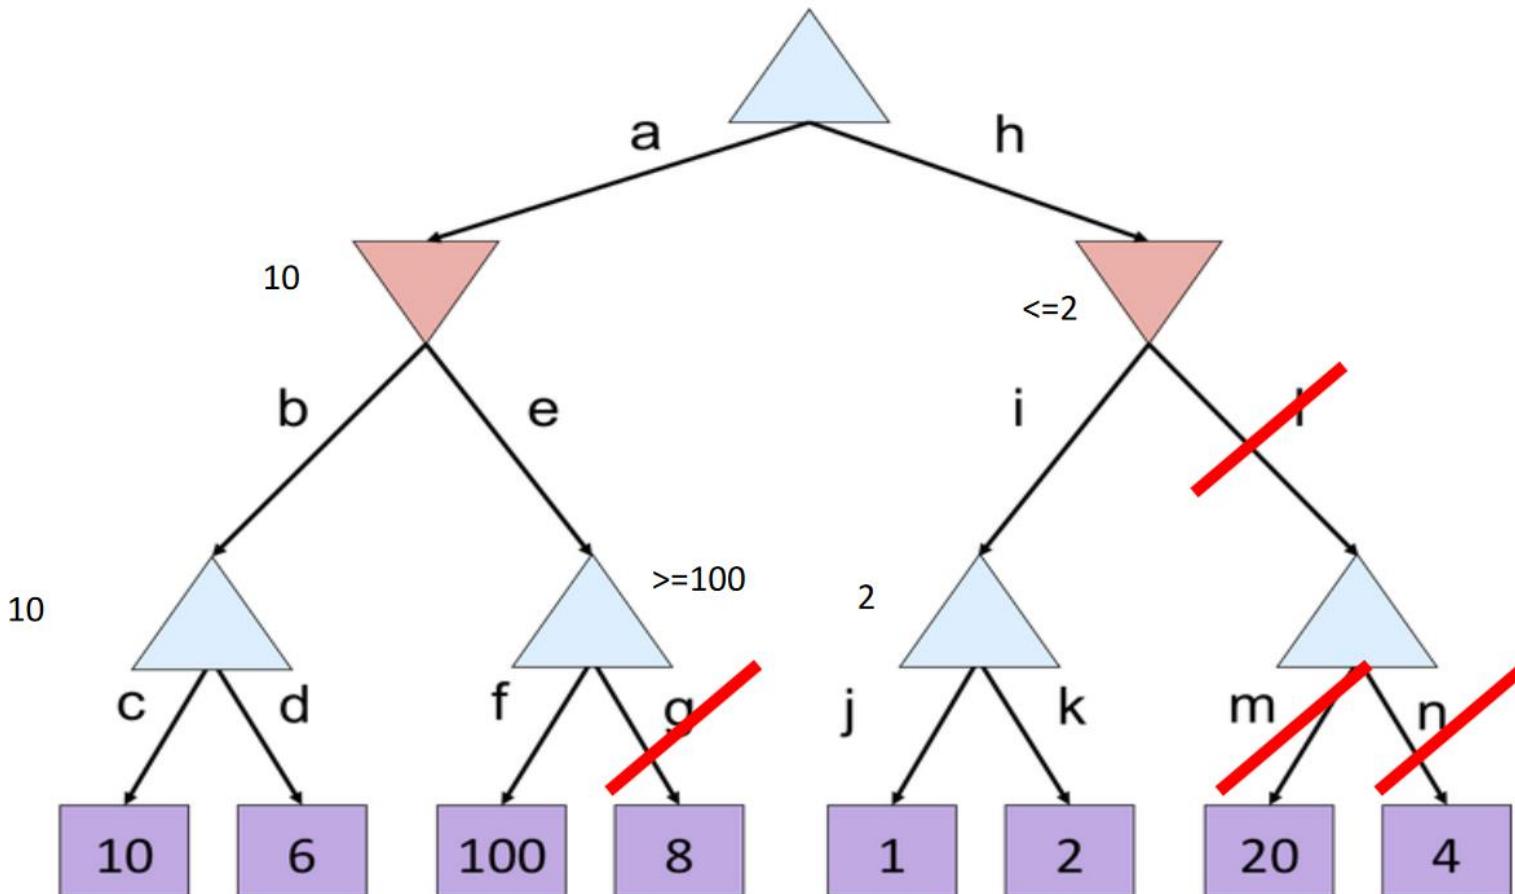
Alpha-Beta Quiz



Alpha-Beta Quiz 2



Alpha-Beta Quiz 2



Alpha-Beta pruning algorithm

```
def minimax_decision(state):
    """
    Returns the optimal action for MAX, using alpha-beta pruning.
    """
    best_value = float("-inf")
    best_action = None
    alpha = float("-inf")
    beta = float("inf")

    for action in actions(state):
        value = minimax_value(result(state, action), alpha, beta)
        if value > best_value:
            best_value = value
            best_action = action

        alpha = max(alpha, best_value) # update alpha

    return best_action
```

Alpha-Beta pruning algorithm

```
def minimax_value(state, alpha, beta):
    """
    Returns the minimax value of a state using alpha-beta
    pruning.
    """

    # Terminal state → return utility
    if terminal_test(state):
        return utility(state)

    # MAX player
    if player(state) == "MAX":
        value = float("-inf")
        for action in actions(state):
            value = max(value, minimax_value(result(state, action),
                                              alpha, beta))
            alpha = max(alpha, value)

        # PRUNING
        if value >= beta:
            return value # prune branch

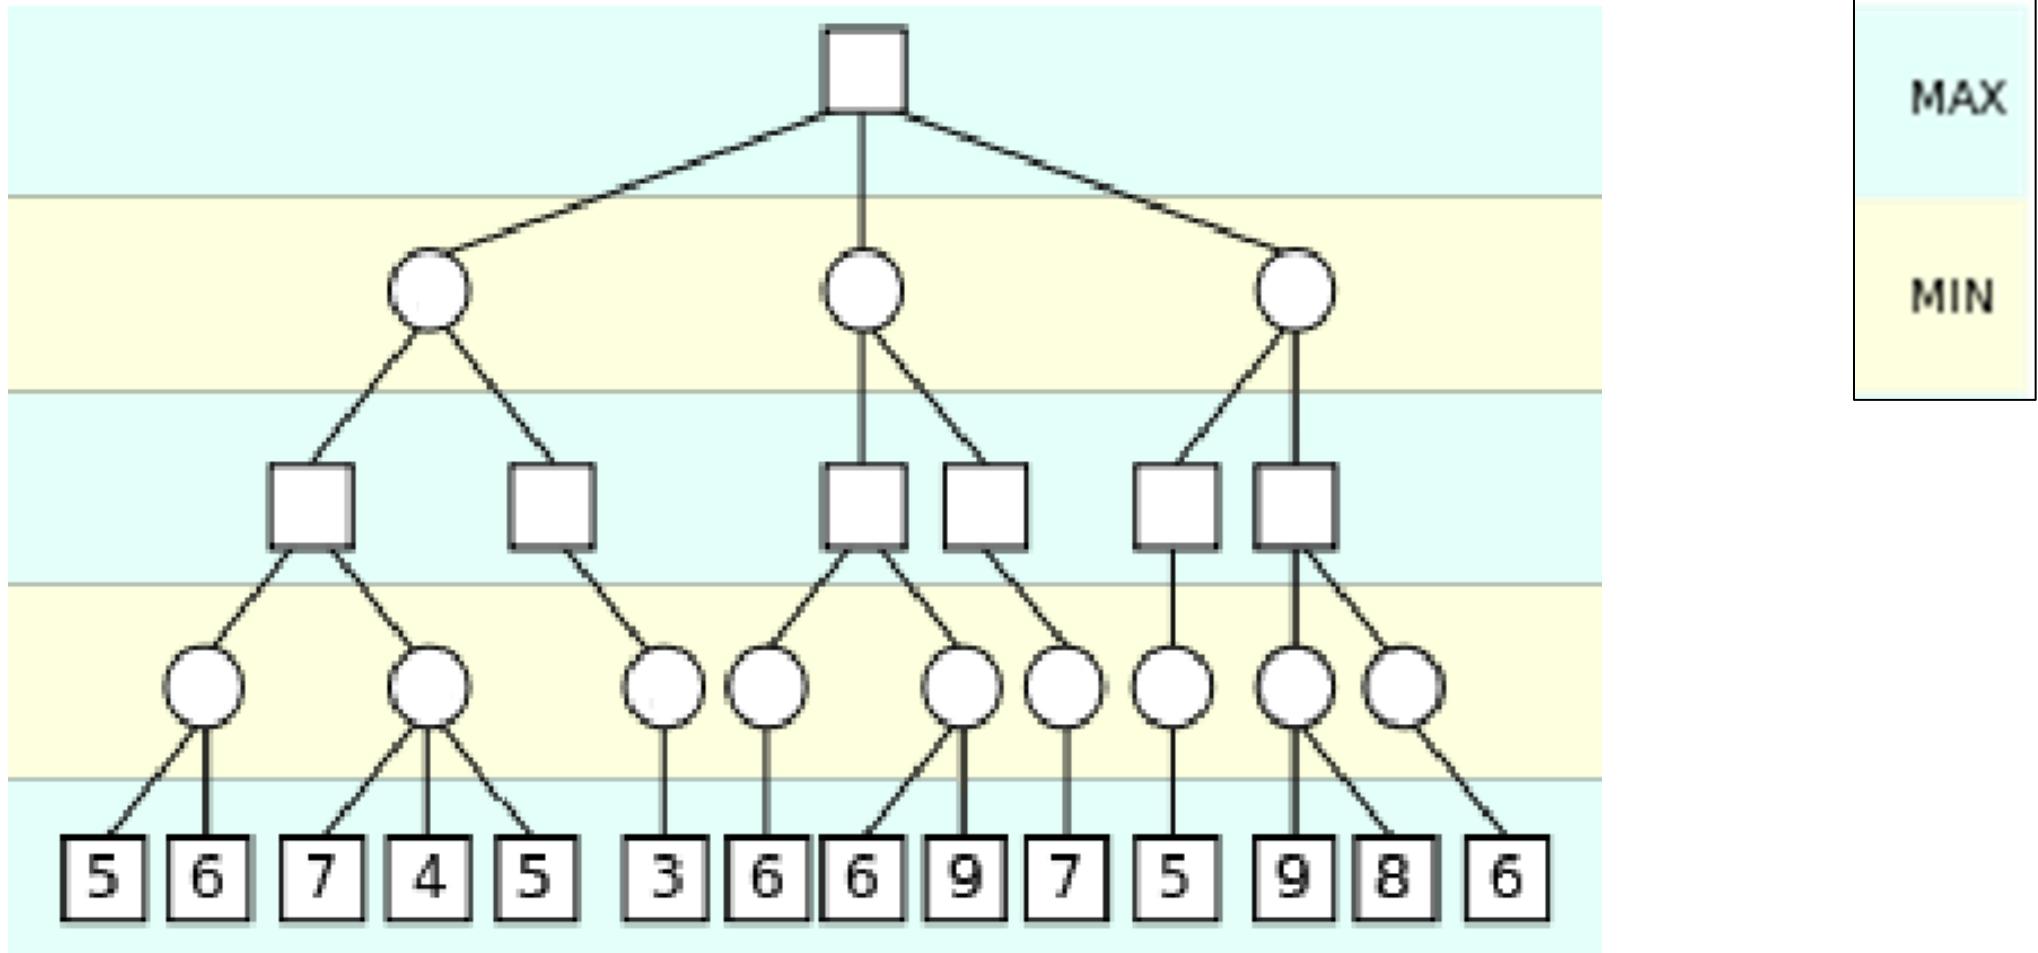
    return value

    # MIN player
else: # player(state) == "MIN"
    value = float("inf")
    for action in actions(state):
        value = min(value, minimax_value(result(state, action),
                                          alpha, beta))
        beta = min(beta, value)

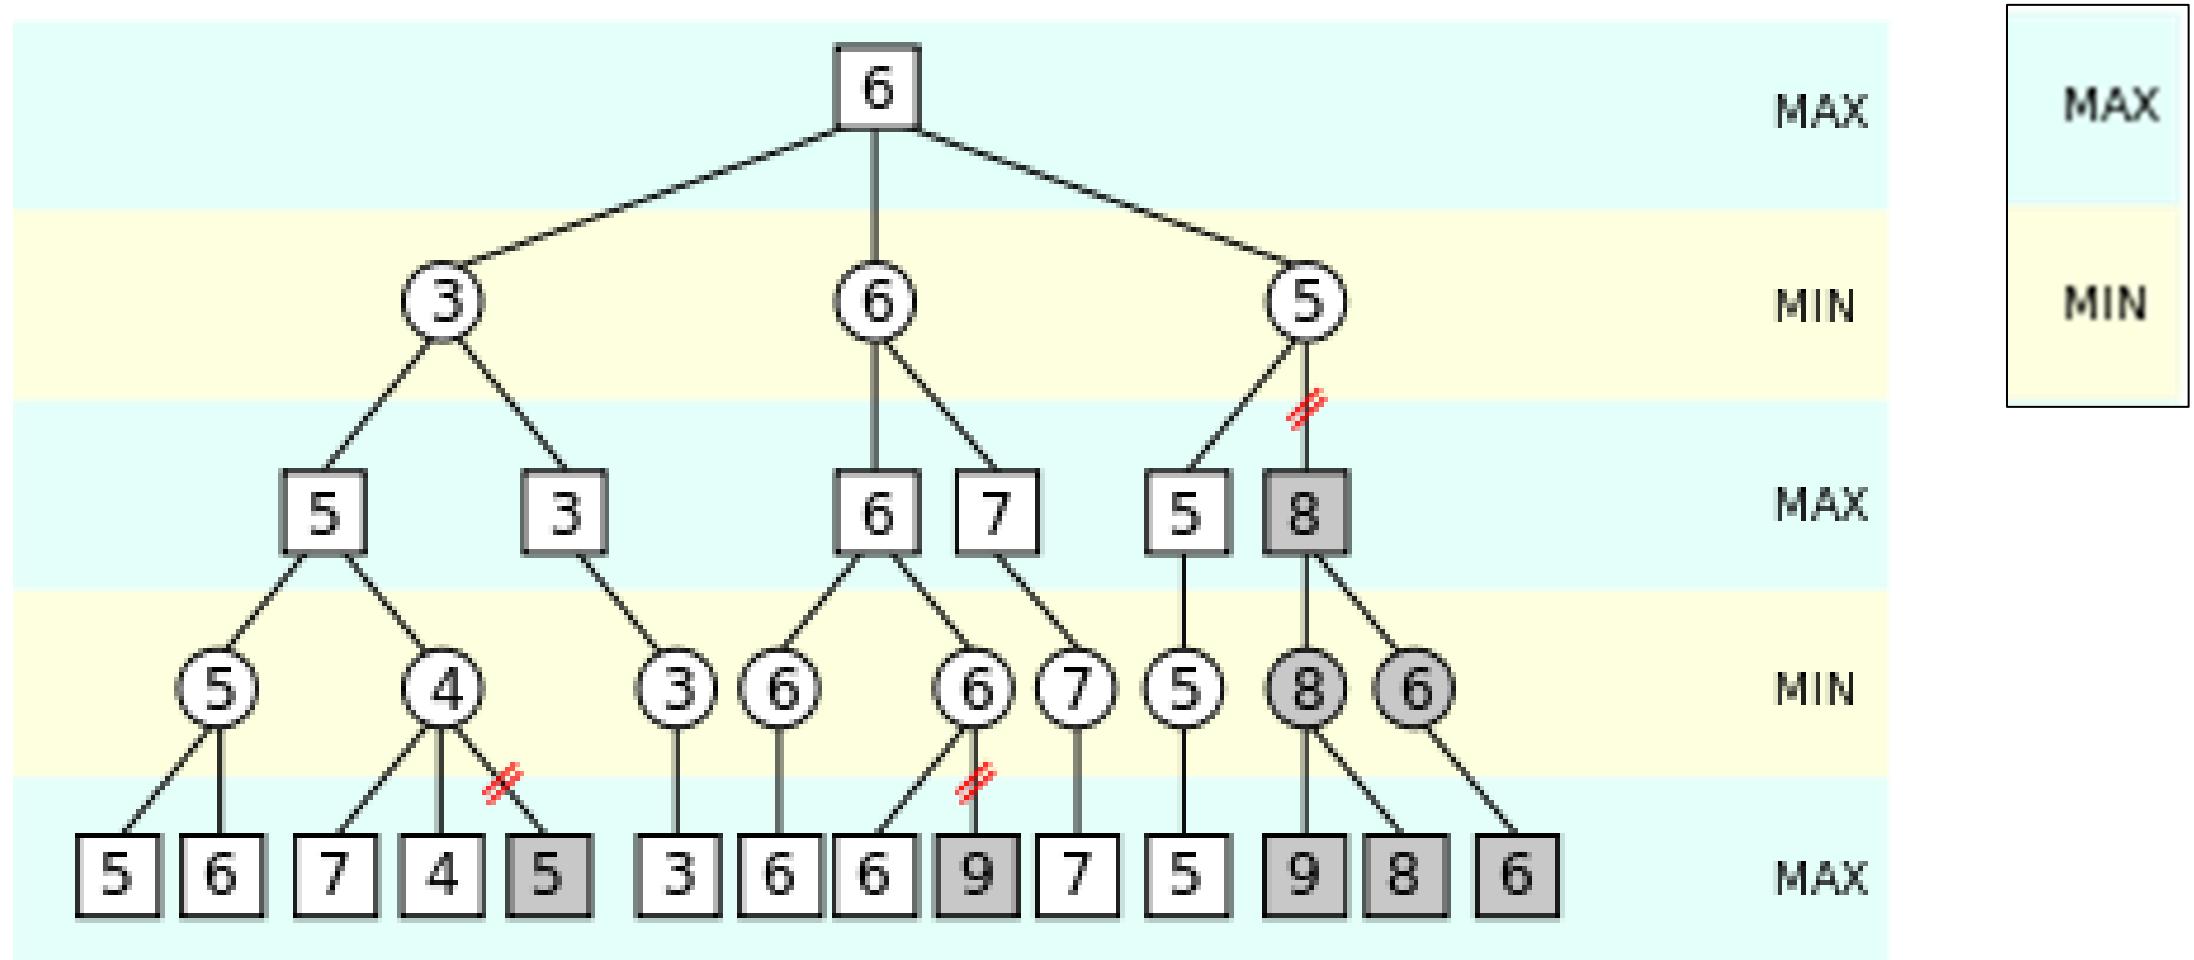
    # PRUNING
    if value <= alpha:
        return value # prune branch

    return value
```

Alpha-Beta Pruning

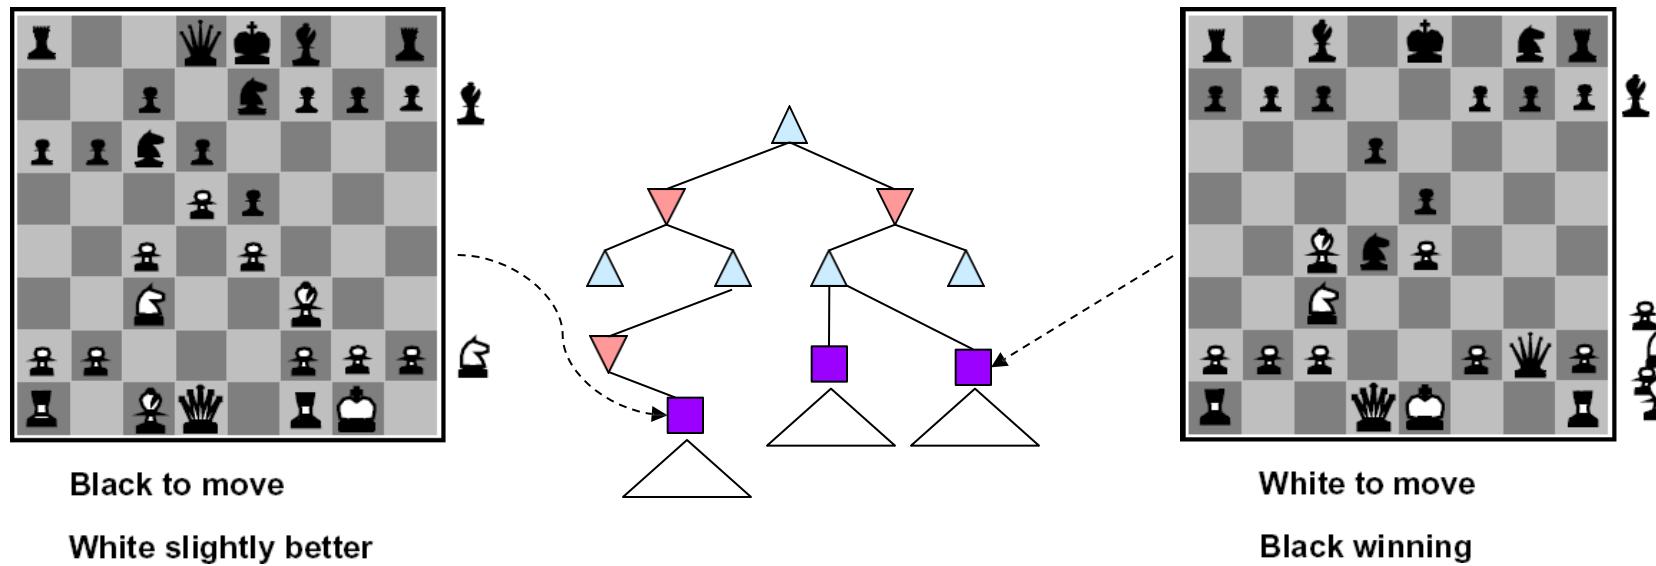


Alpha-Beta Pruning



Evaluation Functions

- Evaluation functions score non-terminals in depth-limited search



- Same idea as heuristics or using mathematical formula

Monte Carlo Tree Search

- Methods based on alpha-beta search assume a fixed horizon
 - Pretty hopeless for Go, with $b > 300$
- MCTS combines two important ideas:
 - ***Evaluation by rollouts*** – play multiple games to termination from a state s (using a simple, fast rollout policy) and count wins and losses
 - ***Selective search*** – explore parts of the tree that will help improve the decision at the root, regardless of depth

MCTS Version 0

- Do N rollouts from each child of the root, record fraction of wins
- Pick the move that gives the best outcome by this metric

